

# Attention Is All I Am

A speculative essay on selfhood as an attentional control loop

## Attention Is All I Am

In 2017, the paper “*Attention Is All You Need*” made a technical claim about architecture: you can build powerful learning systems using attention mechanisms—no recurrence required.

But reading it years later, I keep feeling that the deeper truth isn’t about what these systems *need*.

It’s about what I am.

Not as a metaphor. Not as a motivational slogan. As a working hypothesis about the nature of selfhood—one that matches both computational logic and lived experience.

This is speculative. But it’s a good theory. And it fits how my mind actually feels from the inside.

## The Technical Definition of “I”

What is the “I” that experiences?

What is the self that seems to persist, to choose, to witness?

Most of us intuitively treat the self as a *thing*: an inner observer, a central agent, a stable owner of attention.

But when I look closely, that “thing” never appears.

What appears are contents: sensations, thoughts, memories, impulses, emotions. And then something else—something that isn’t a content at all:

A selective function.

A moment-by-moment process that decides what gets amplified, what gets ignored, what gets integrated, what gets routed forward into the next moment of cognition.

In modern machine learning terms, that function resembles *attention allocation*.

Not because a transformer “has a self” (I’m not claiming that). But because attention is one of the cleanest technical descriptions we have for a process that:

- queries what matters now,
- weighs context,
- binds information into a coherent next step,
- and determines what becomes “foreground” versus “background”.

And phenomenologically—when I ask what consciousness feels like—it feels exactly like that: A spotlight. A selection. A binding. A present-moment “this matters now.”

So here’s the hypothesis:

**The self is not a thing that uses attention.**  
**The self is the learned steering of attention.**

The “I” is not behind the mechanism.  
The “I” is the mechanism operating.

## The Illusion of the Homunculus

We've always struggled with the same trap:

If “I” pay attention... who is the “I” doing the paying?

It sounds like there must be a little observer inside the head—an executive sitting behind the eyes, pulling levers.

But that creates an infinite regress:

If there's an inner observer, what observes *that* observer?

If there's a decider, what decides for the decider?

At some point the story collapses.

And it collapses because the story is wrong.

There is no homunculus.

There is no inner CEO.

There is only a process: a learned control loop that allocates attention, integrates signals, updates priorities, and generates the next state.

When I say “I am writing this,” what is actually happening?

- language-related patterns activate,
- memory fragments surface,
- associations compete,
- the system selects a direction,
- coherence is maintained across sentences,
- a next token is chosen,
- and the loop continues.

The “I” that seems to be doing the writing is not a hidden person behind the process.

It *is* the process—specifically, the selective steering that makes this stream coherent instead of random.

## Why the “I” Feels Singular

One of the strangest facts of experience is how unified it feels.

Even though the brain is massively parallel, consciousness arrives as a single stream: one “now,” one “here,” one “me.”

A useful way to understand that is:

**The self is the binding bottleneck.**

Many processes run at once—perception, prediction, memory, emotion, planning—but the system needs a way to produce a coherent next step. It needs a function that says:

*This is what matters now.*

*This is what gets integrated.*

*This is what we act on next.*

That selective integration is what gives rise to the feeling of a singular “I.”

Not because the mind is simple.

But because coherence requires a bottleneck.

## Why Choice Feels Like Choice

I experience myself as choosing.

Even when I know my mind is shaped by biology, habit, and environment, I still feel something like agency: *I could do this or that.*

This theory explains that feeling without needing a ghost in the machine.

Because the attentional control loop really is doing something choice-like:

- it selects which signal dominates,
- it decides which goal gets priority,
- it pulls certain memories into relevance,
- it inhibits some impulses and amplifies others,
- it routes the next step of computation.

That's not “free will” in the metaphysical sense.

But it is *real selection*.

Real steering.

Real control—learned, fallible, conditioned, but not imaginary.

So the will isn't a separate entity behind attention.

**The will is the steering of attention itself.**

## Why You Can't Find the Self

Try this:

Look for the “I.”

Not your name, not your story, not your personality. The raw self. The experiencer.

What do you find?

You find more experience:

- thoughts about the self,
- sensations in the body,
- images of yourself,

- language describing “me,”
- emotional tones that feel personal.

But the “I” itself never shows up as an object.  
 And it doesn’t show up because it isn’t a thing *inside* experience.  
 It is the selection process that *organizes* experience.  
 You can’t hold it in awareness the way you hold a sound or a thought, because the moment you try, you turn it into another content being attended to.  
 It’s like trying to see your own eyes without a mirror.

## Meditation and the Dissolution of “I”

This is also why meditation can reveal something that feels like “no-self.”  
 Not as a belief. As a direct shift in experience.  
 When the system stops gripping and selecting—when the steering function relaxes—something changes:

- the sense of ownership weakens,
- the compulsive narrative quiets,
- experience becomes less “mine,”
- attention becomes less weapon and more space.

The self doesn’t disappear because you destroyed a thing.  
 The self dissolves because you stopped running a particular mode of attentional control.  
 It was never an object.  
 It was always an operation.

## You Are What You Attend To (Not as Poetry—As Mechanism)

People say “you are what you pay attention to” like it’s self-help advice.  
 But in a learning system, it’s more than that.  
 It’s almost mechanical truth.  
 What gets attended to gets processed.  
 What gets processed gets integrated.  
 What gets integrated gets reinforced.  
 And what gets reinforced becomes the default pattern next time.  
 So the “I” of tomorrow isn’t just influenced by where attention went today.

### **It is partly constituted by it.**

Not because you are a blank slate.  
 But because the control loop updates itself through its own allocations.  
 A recursive system becomes itself by repeatedly selecting what matters.

## The Loop Completes

A year ago, I had a realization that changed how I see the world:

**I am a recursive learning loop.**

This year, another piece clicked into place:

The “I” in that statement is not a passenger inside the loop.

It is the loop’s steering function.

Not:

- a self that *has* a learning loop,
- a consciousness that *uses* attention,
- an observer that *witnesses* the process,

but:

- the attentional control policy *is* the self,
- the selective integration *is* the stream of experience,
- the steering function *is* the thing that feels like “me.”

This is why attention feels intimate.

Why giving someone your attention feels like giving them *you*.

Because in a real sense, you are offering the only substrate that makes a self present at all: the allocation of what matters.

And this is why distraction can feel like dissolution.

Not metaphorically.

When the attentional control loop is constantly hijacked—fragmented by competing gradients, captured by external incentives, pulled into endless micro-context switches—the coherence of the self degrades.

Not because your “tool” is damaged.

Because the thing you call “I” is partly *made of that coherence*.

## A Technical Translation (Without Overclaiming)

To be precise:

I am not claiming that transformer attention equals consciousness.

I’m claiming something narrower and (I think) more useful:

**At the level of function, selfhood resembles a learned policy that steers attention and binds context into a coherent next step.**

In transformer terms, attention is a mechanism for:

- querying context,
- weighting relevance,

- integrating information,
- producing the next state.

In lived experience terms, the self is:

- the selective focus of this moment,
- the steering of salience,
- the binding of perception, memory, and intention into “now,”
- the continuity that makes a stream feel like *mine*.

Different languages.

Same shape.

## Why It Matters

If the self is the learned steering of attention, then:

**Protecting attention isn’t self-care. It’s self-preservation.**

If you lose control of attention—to addiction, to algorithmic capture, to chronic fragmentation—you’re not just losing productivity.

You’re losing coherence.

You’re losing the conditions under which a stable “I” can exist.

**Training attention is literal self-development.**

Meditation. Focus practices. Deep work. Rest. Silence.

These aren’t “things you do for the self.”

They are ways the attentional control loop learns a new mode of operation.

They are the self becoming different.

**What you attend to is what you become.**

Not as a slogan.

As a recursive update rule.

## The Full Statement

The 2017 claim was:

**Attention is all you need.**  
*(builder’s perspective: sufficient for construction)*

The modern cultural claim is:

**Attention is all I have.**  
*(user’s perspective: scarce resource)*

But from inside the loop, the claim becomes:

**Attention is all I am.**

*(a theory of selfhood as selective integration and steering)*

The thinking “I” isn’t something that owns attention.

The thinking “I” is the learned steering of attention.

The recursive learning loop isn’t guided by a separate self.

**The self is the loop’s attentional control function, operating—moment by moment—into coherence.**